

COL216

Computer Architecture

Memory Organization

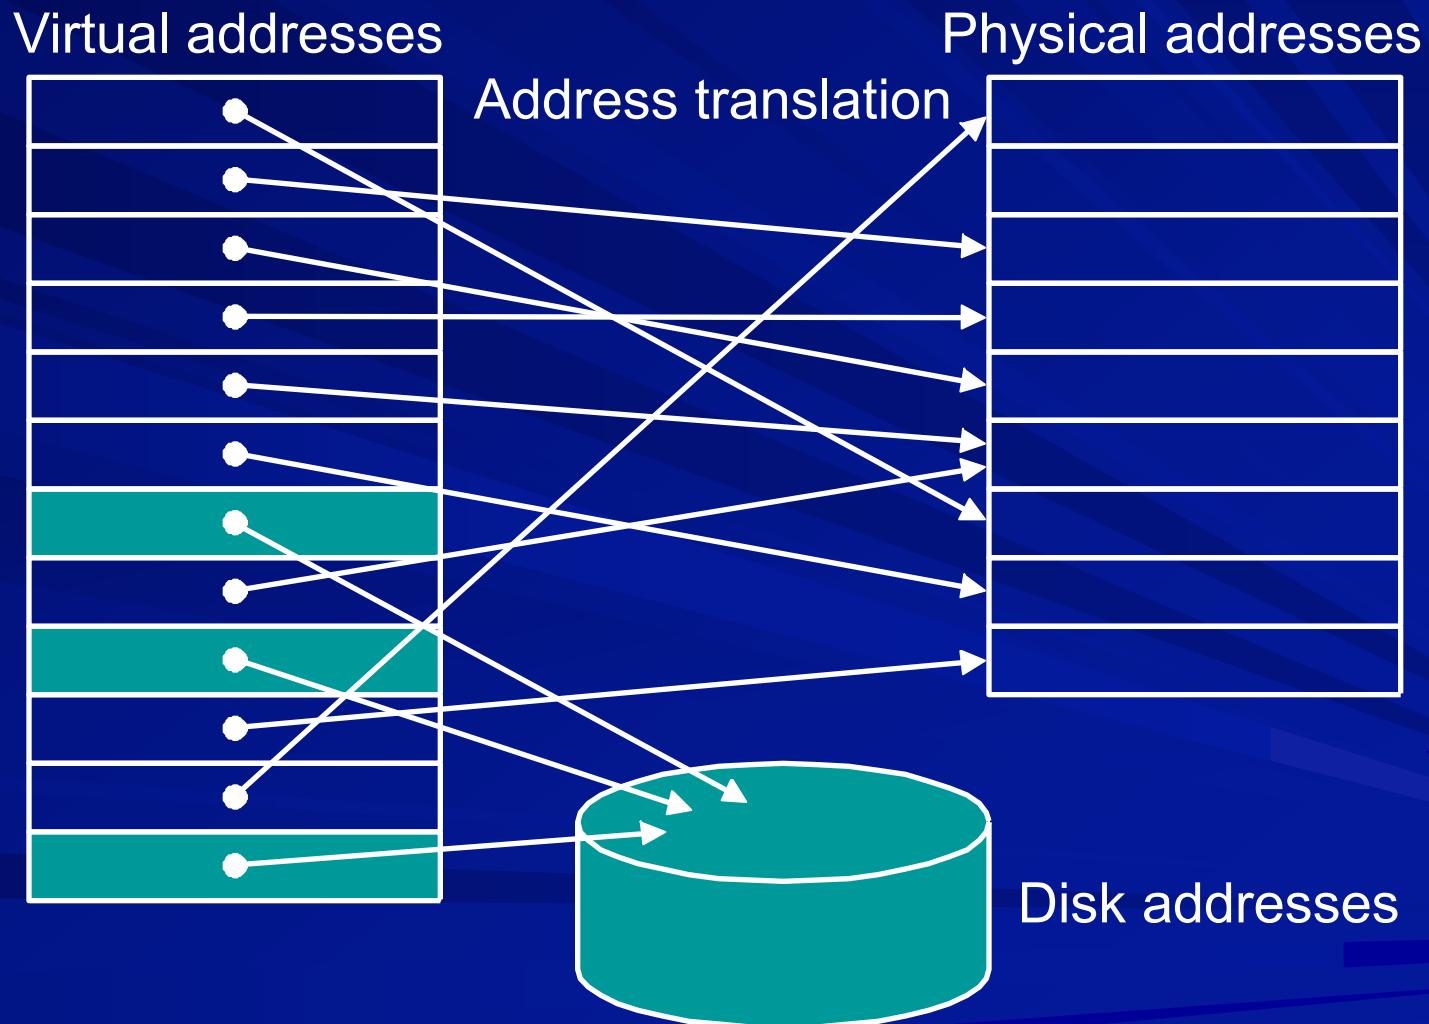
Virtual Memory

12th March 2022

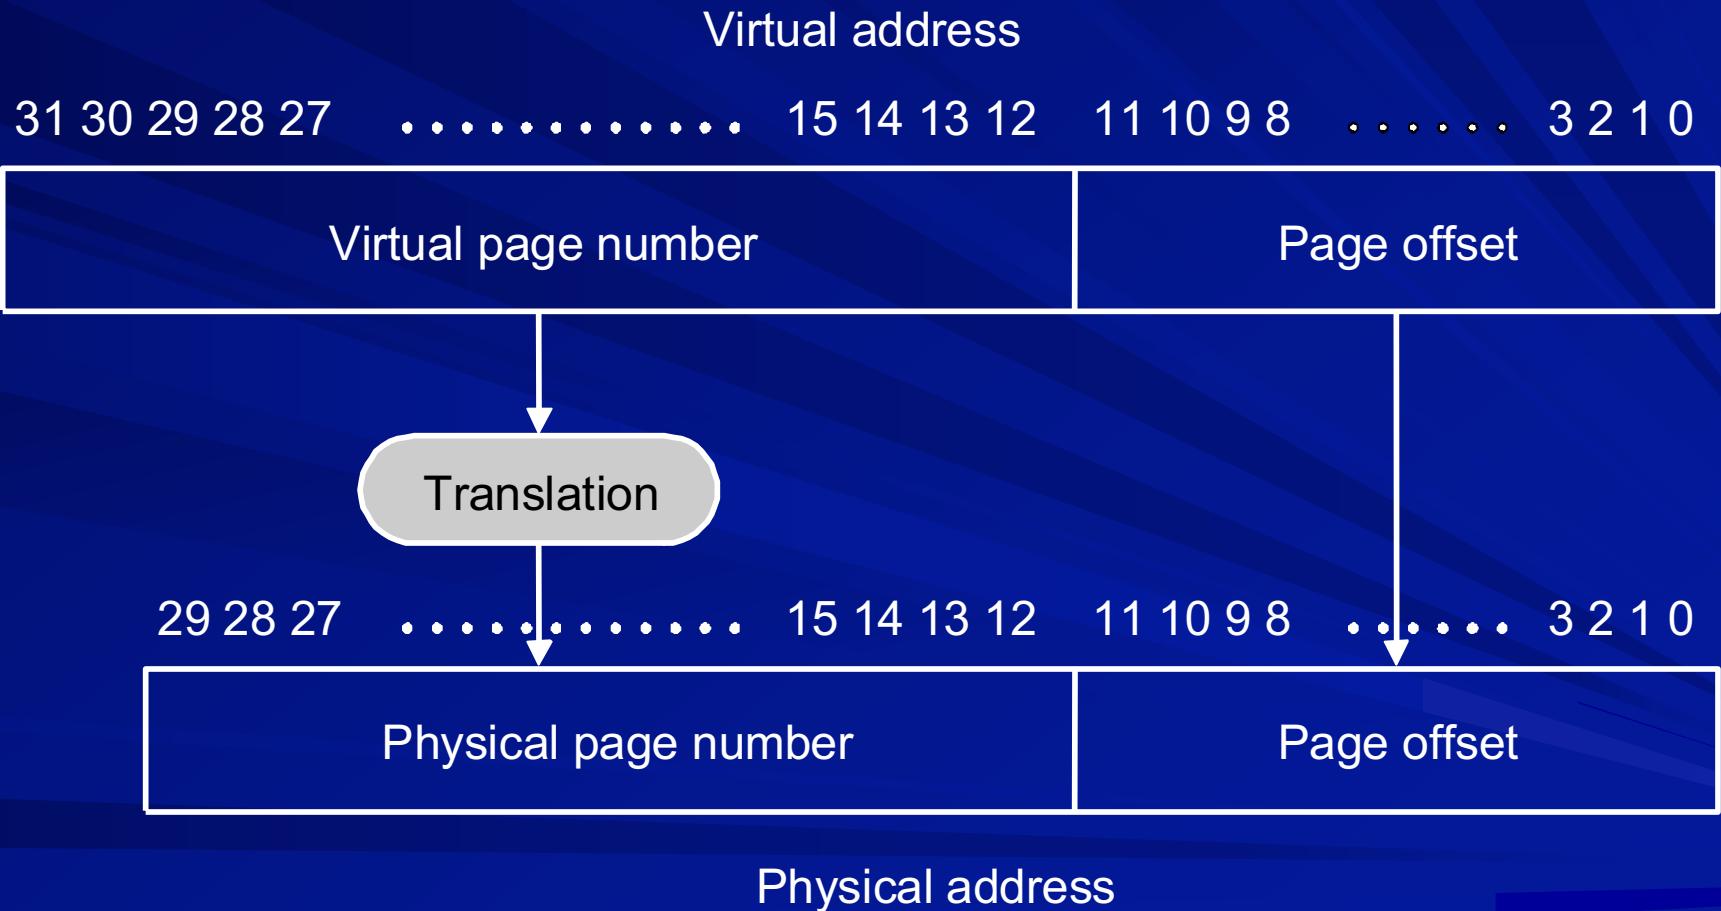
Virtual memory objectives

- To overcome the limitation of the physical memory size
- To allow multiple programs to share memory
 - Protection
 - Program relocation

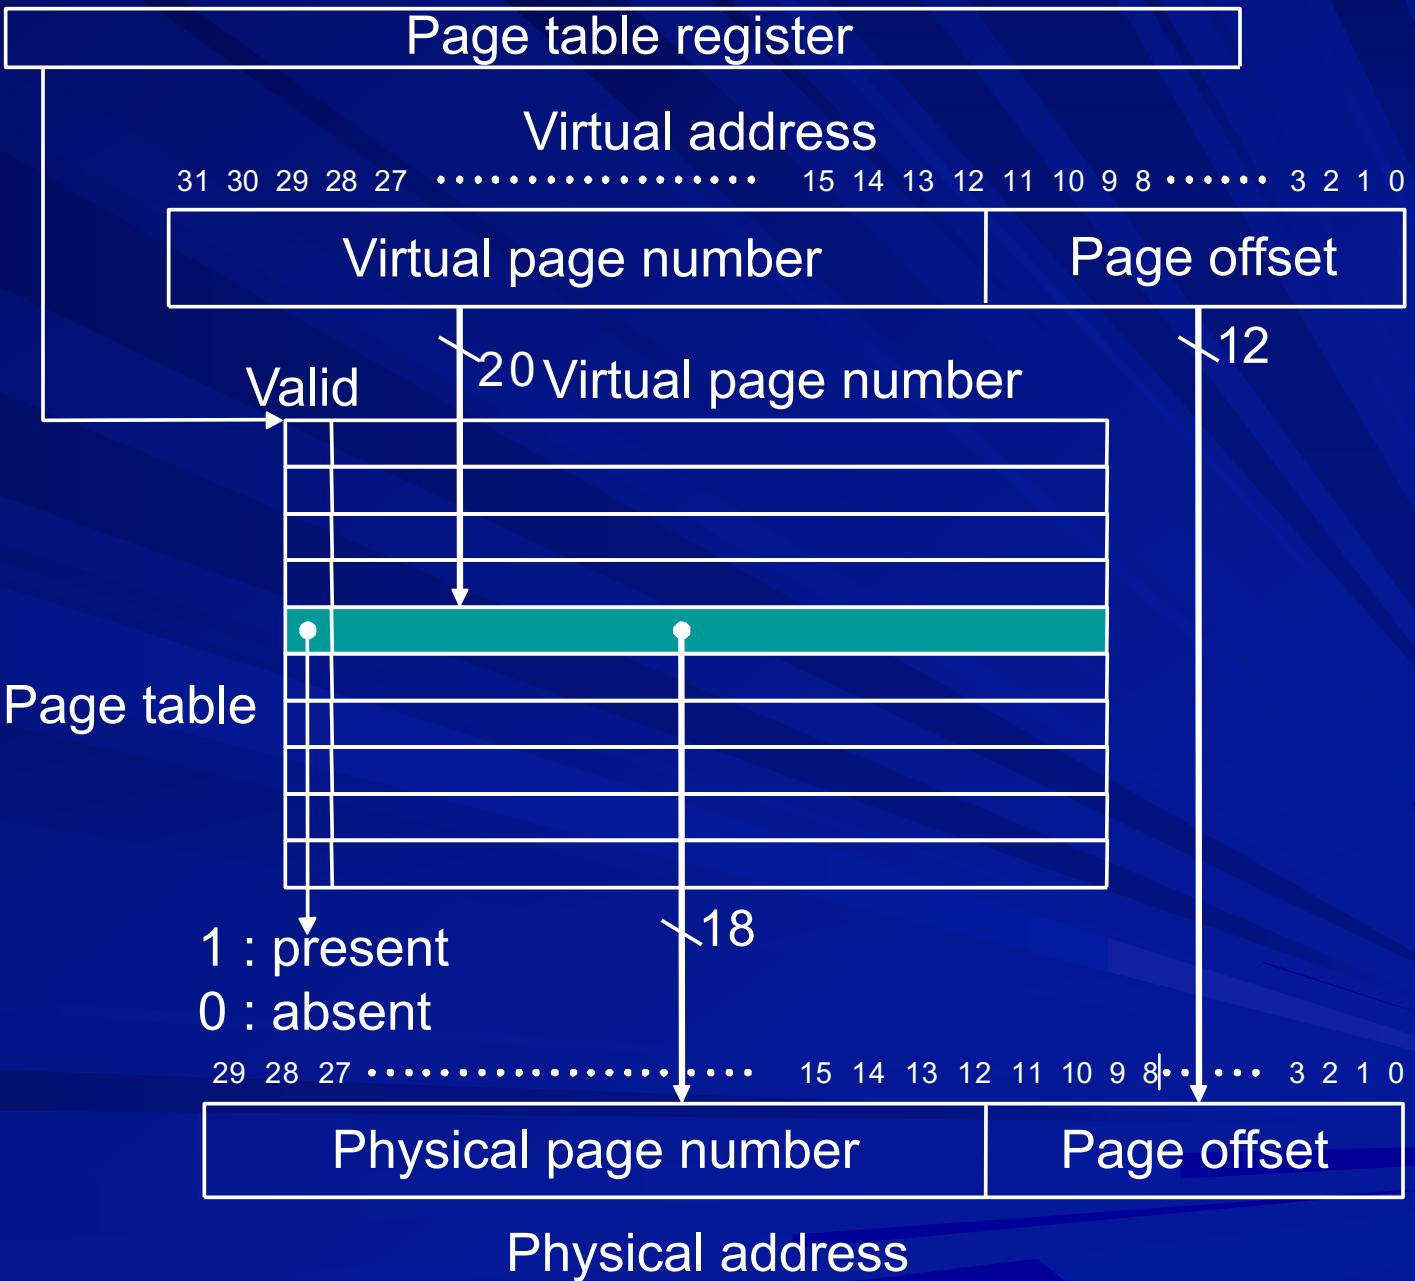
Virtual to physical mapping



Address translation



Page table



Mapping with page table

virtual page number

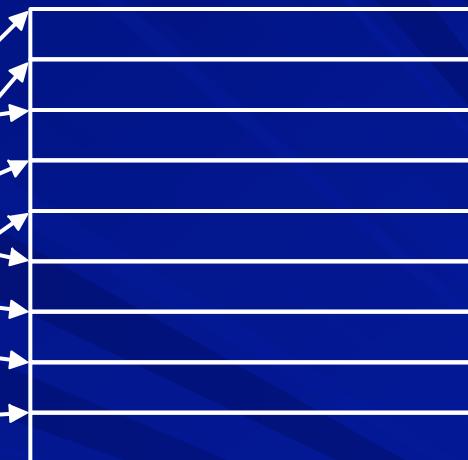


Page table:

v physical page
or disk addr

1	•
1	•
1	•
1	•
0	•
1	•
1	•
0	•
1	•
1	•
0	•
1	•

Physical memory



Disk storage



Virtual memory vs cache : similar but important differences

- Speed difference
 - cache and main memory : one order of magnitude
 - main memory and HDD : several orders of magnitude
- Response to a miss
 - cache : must be handled by h/w, can't afford to switch context
 - VM : can't keep CPU waiting, must switch context, can be conveniently handled by s/w
- Terminology differences
 - page vs block(cache line), page fault vs miss, page table vs cache directory etc

Differences contd.

■ Miss rate

- VM can only afford extremely low miss rates
- main memory much larger than cache, implies much lower miss rate

■ Policies

- pages much larger than blocks (~4KB - 16 KB) : capture large spatial locality, amortize transfer time
- fully flexible mapping (assoc mem not used)
- always write back
- good approximation to LRU

Back to page table

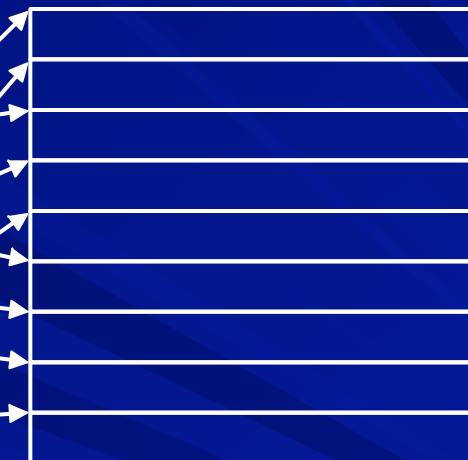
virtual page number



Page table:
physical page
or disk addr

1	•
1	•
1	•
1	•
0	•
1	•
1	•
0	•
1	•
1	•
0	•
1	•

Physical memory



Disk storage



Where is page table stored?

How big is it?

- Suppose virtual address = 32 bits, page size = 4 KB, page table entry = 4 B
- then number of page table entries = number of pages = $2^{32} / 2^{12} = 2^{20}$
- page table size = $2^{20} \times 2^2 = 4 \text{ MB}$
- hundreds of processes

Store in dedicated memory? main memory?

Handling large page tables

- Bound the page table size
- Exploit sparseness
- Use multiple levels
- Page the page table

Bounding the page table size

- Keep a bound register
 - Allocate as much space as necessary
 - Address space expands in one direction
- Split the virtual space into two parts, each with its own bound, e.g. heap + stack
 - Two segments can grow independently
- Sparseness not exploited

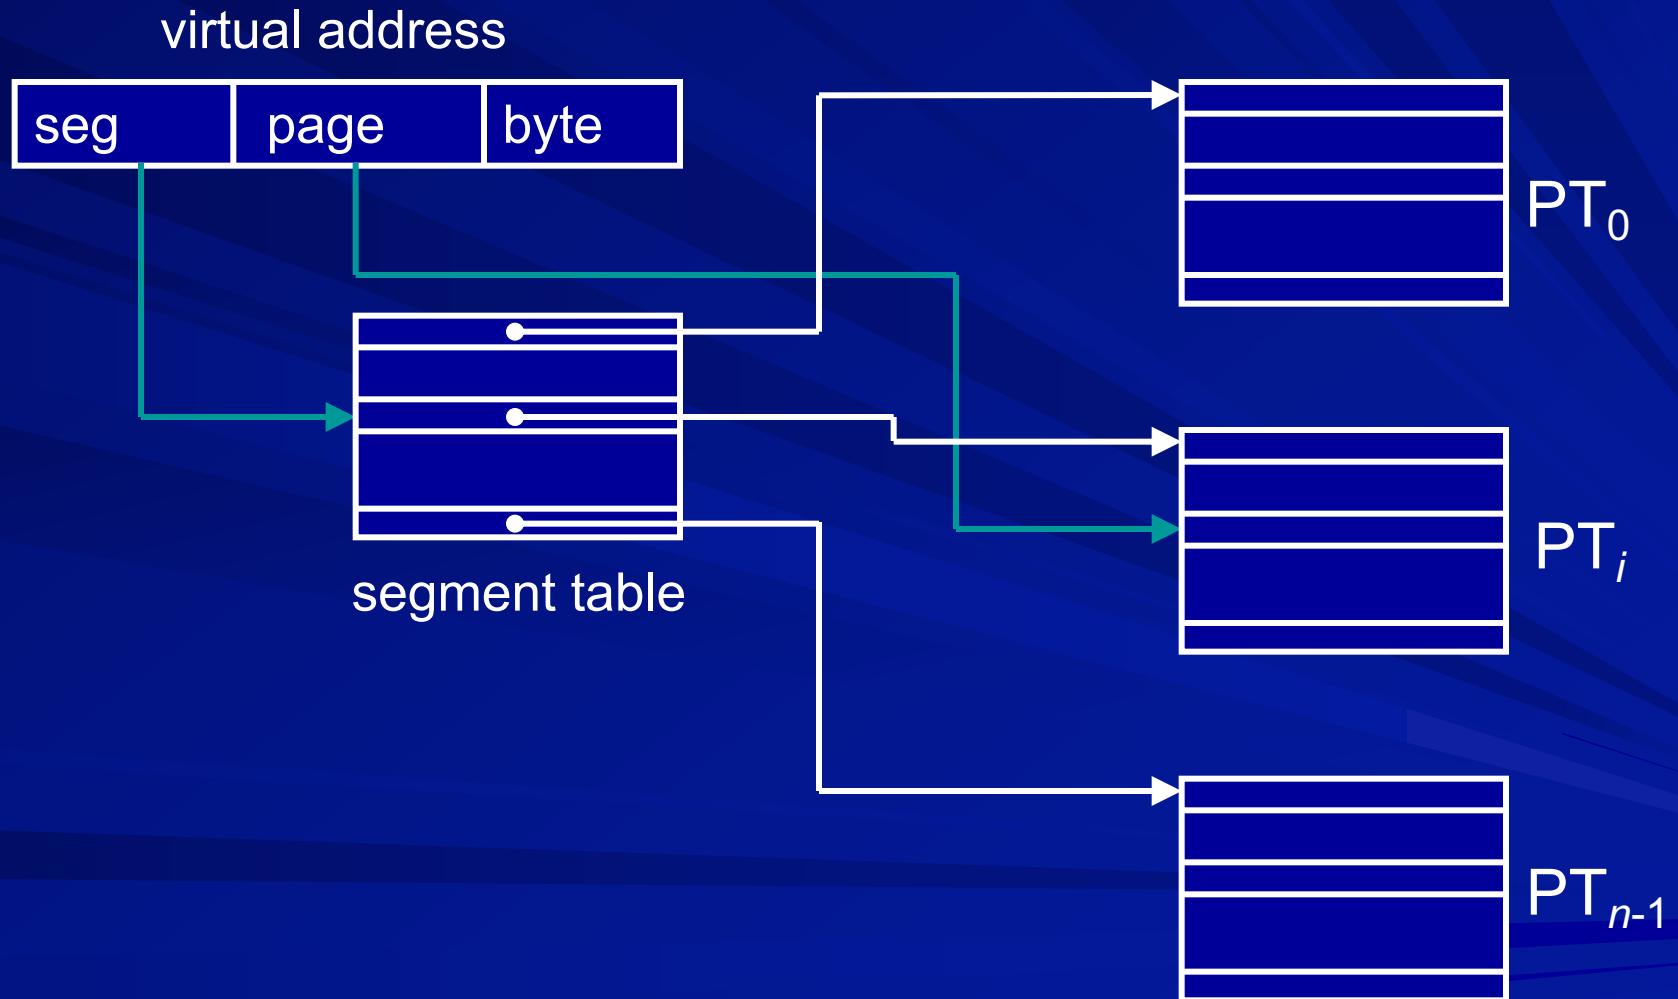
Exploit sparseness

- Page table: one entry per virtual page,
Cache: one entry per cache block
 - Associative memory impractical for page table
- Hashing can be used - inverted page table
 - Access is more complex than simply indexing into the page table

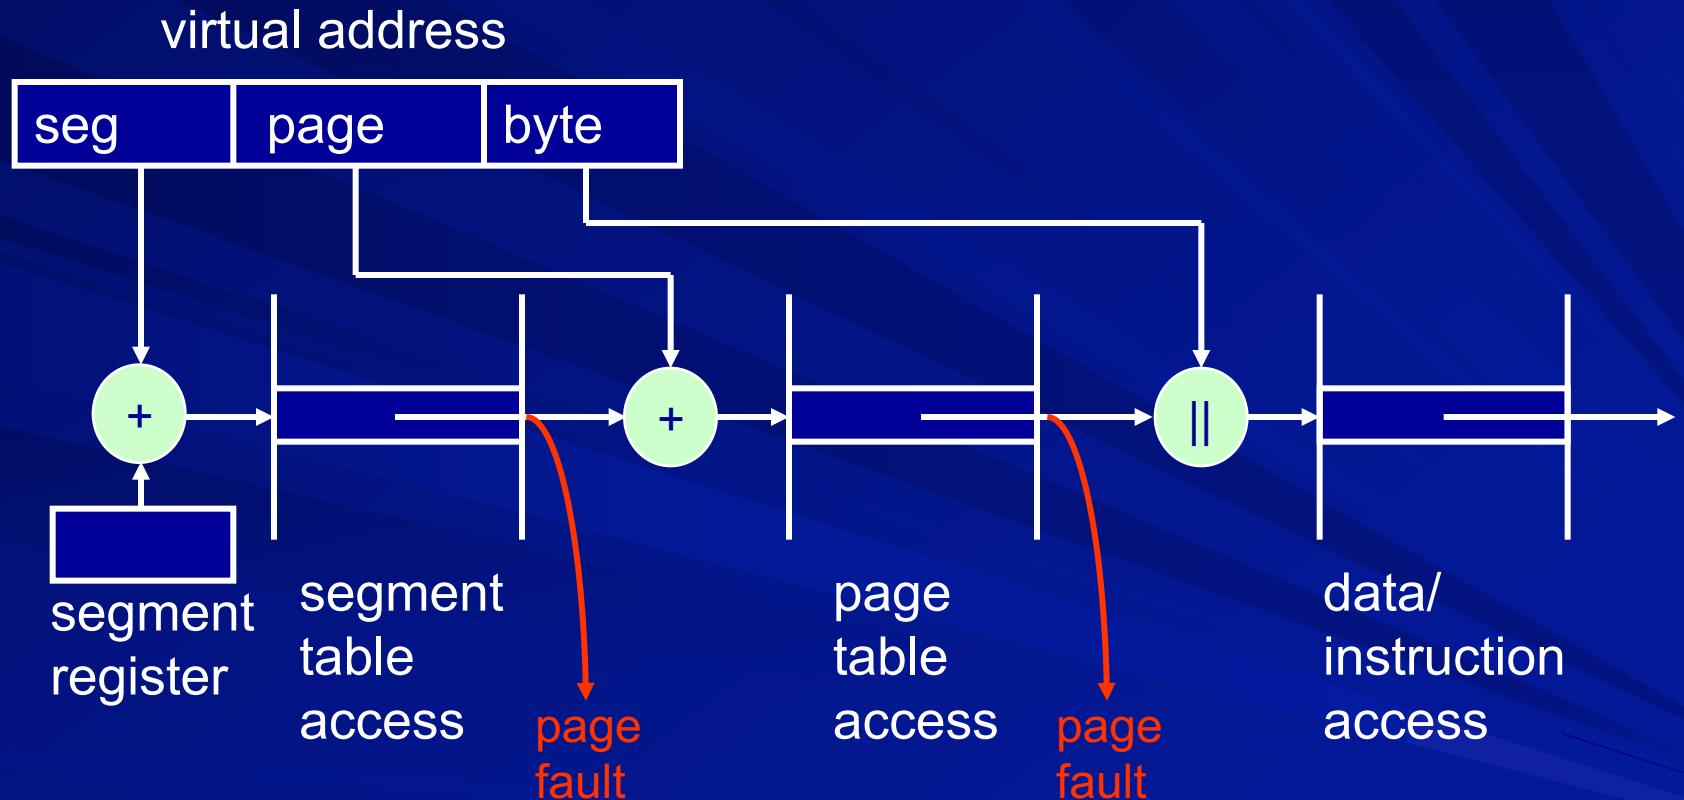
Two level page tables

- Virtual space divided into segments
- A smaller page table for each segment
- All page tables need not be in physical memory
- A segment table keeps track of where the page tables of different segments are located

2 Level page table structure



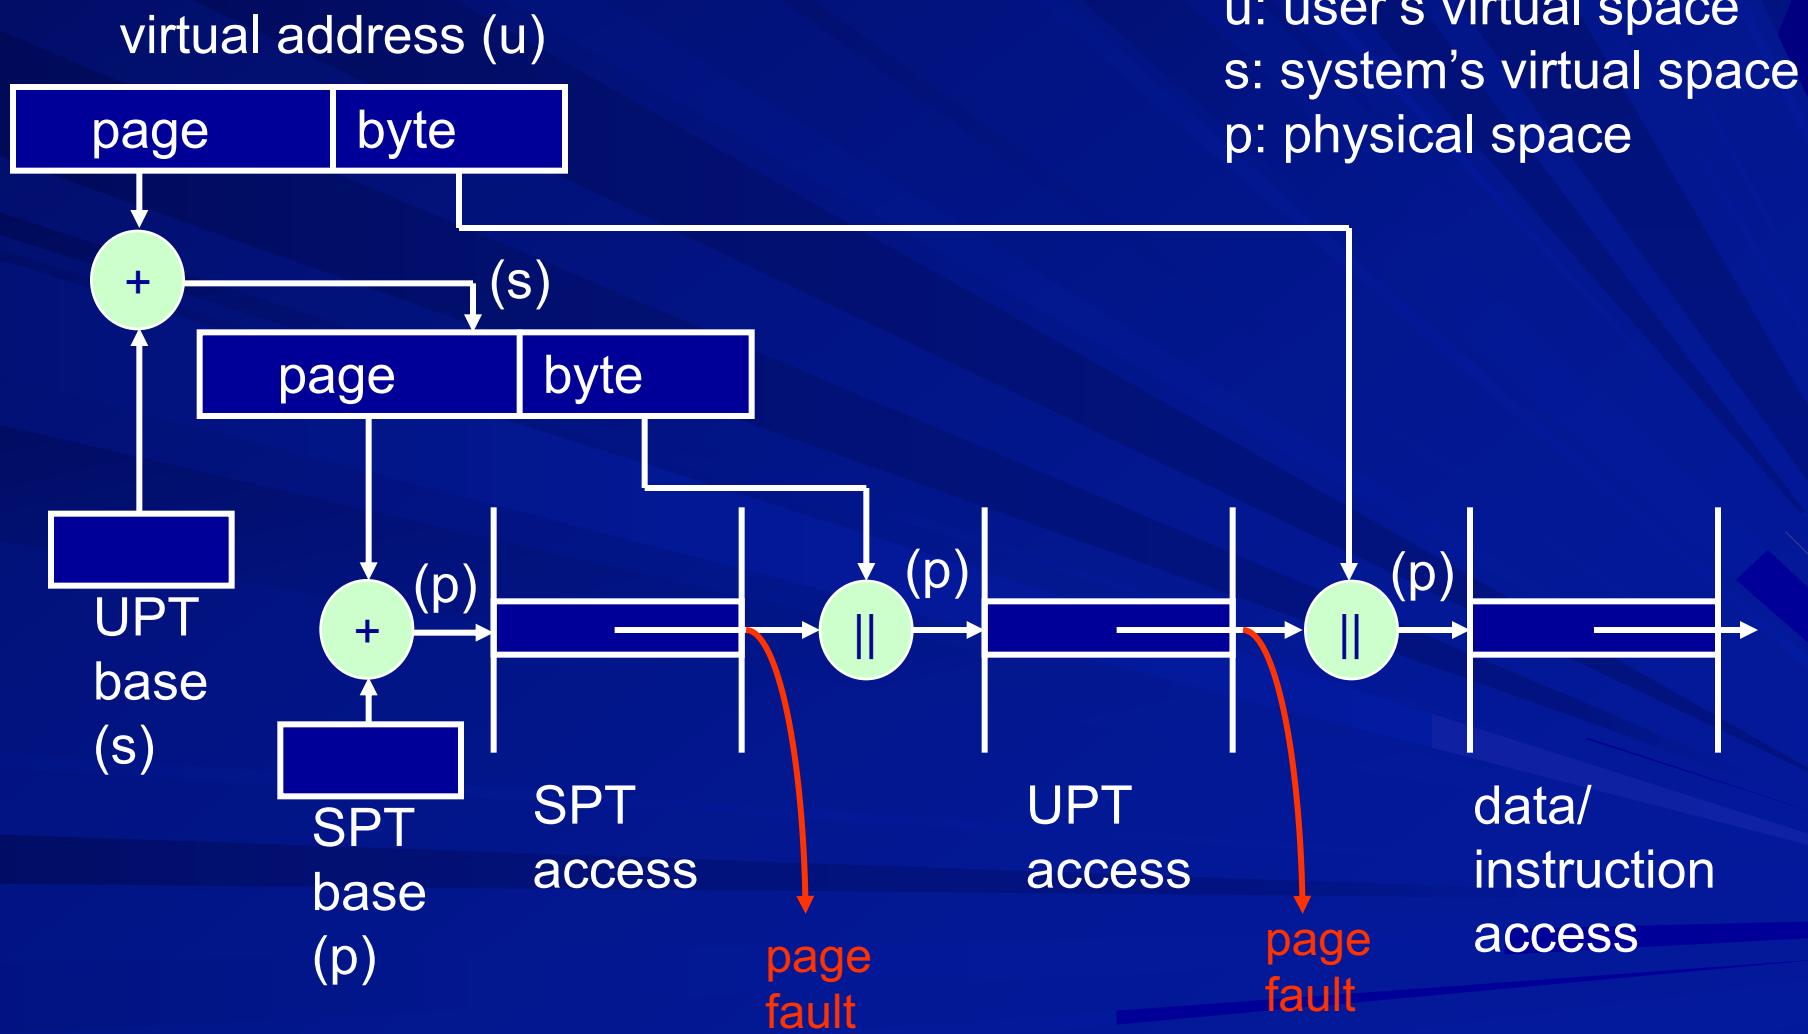
Memory access with 2 level page table



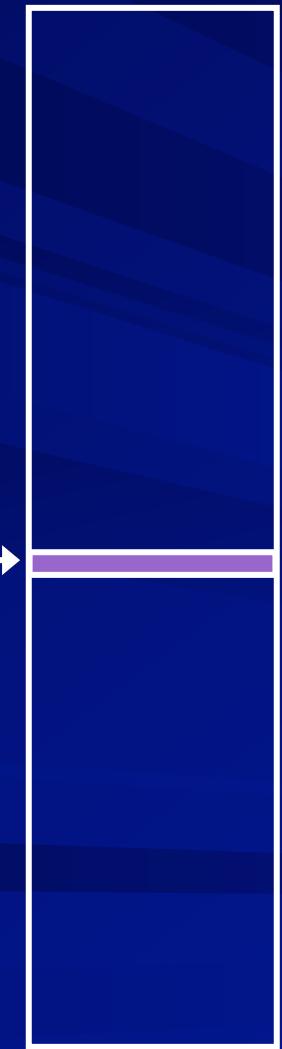
Keep page table in virtual space

- Page table is paged!
- Only active pages of page table are brought into physical memory
- How do you locate these pages? - need another page table!
- User page tables are kept in system's virtual space
- System's page table (or at least a part of it) must be kept in physical memory and accessed directly

Memory access with paged page table



user
virtual space



physical
space



user
virtual space

X →



physical
space

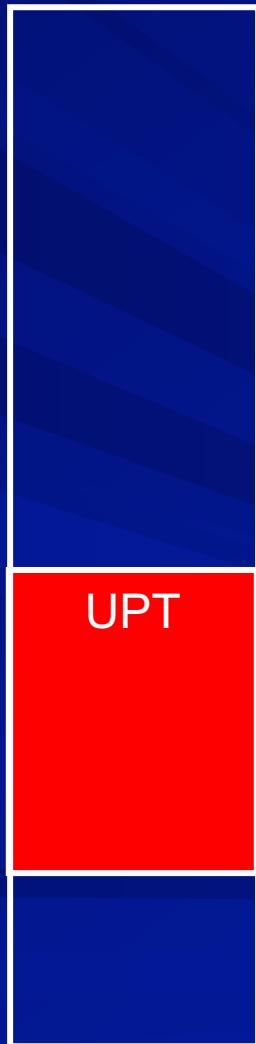


physical
page

user
virtual space



system
virtual space

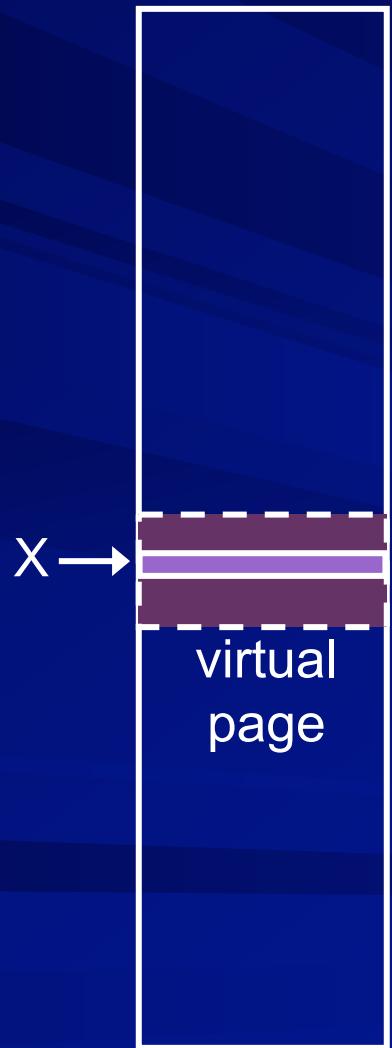


physical
space

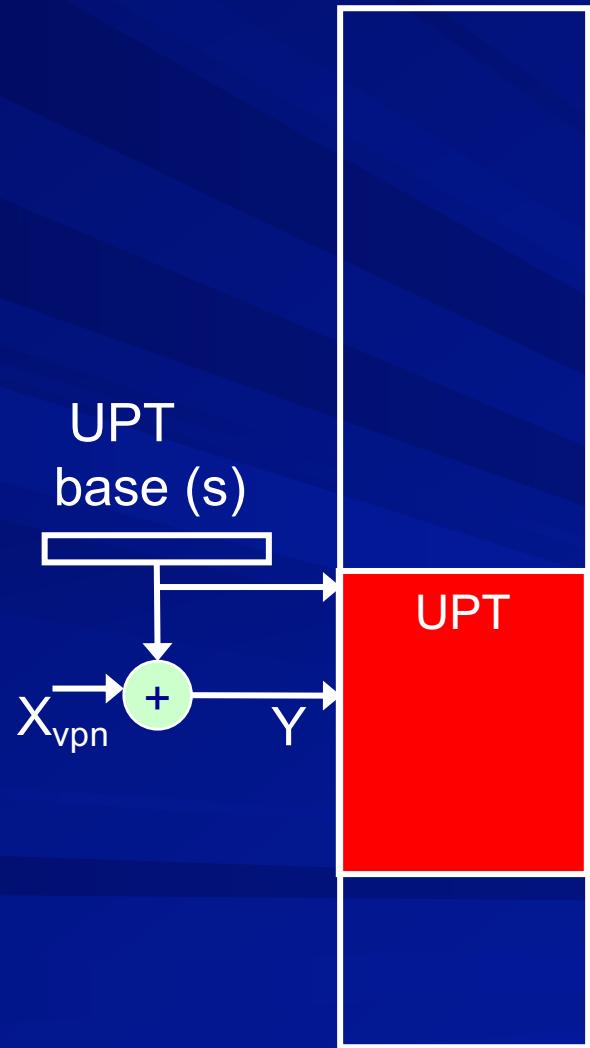


physical
page

user
virtual space



system
virtual space



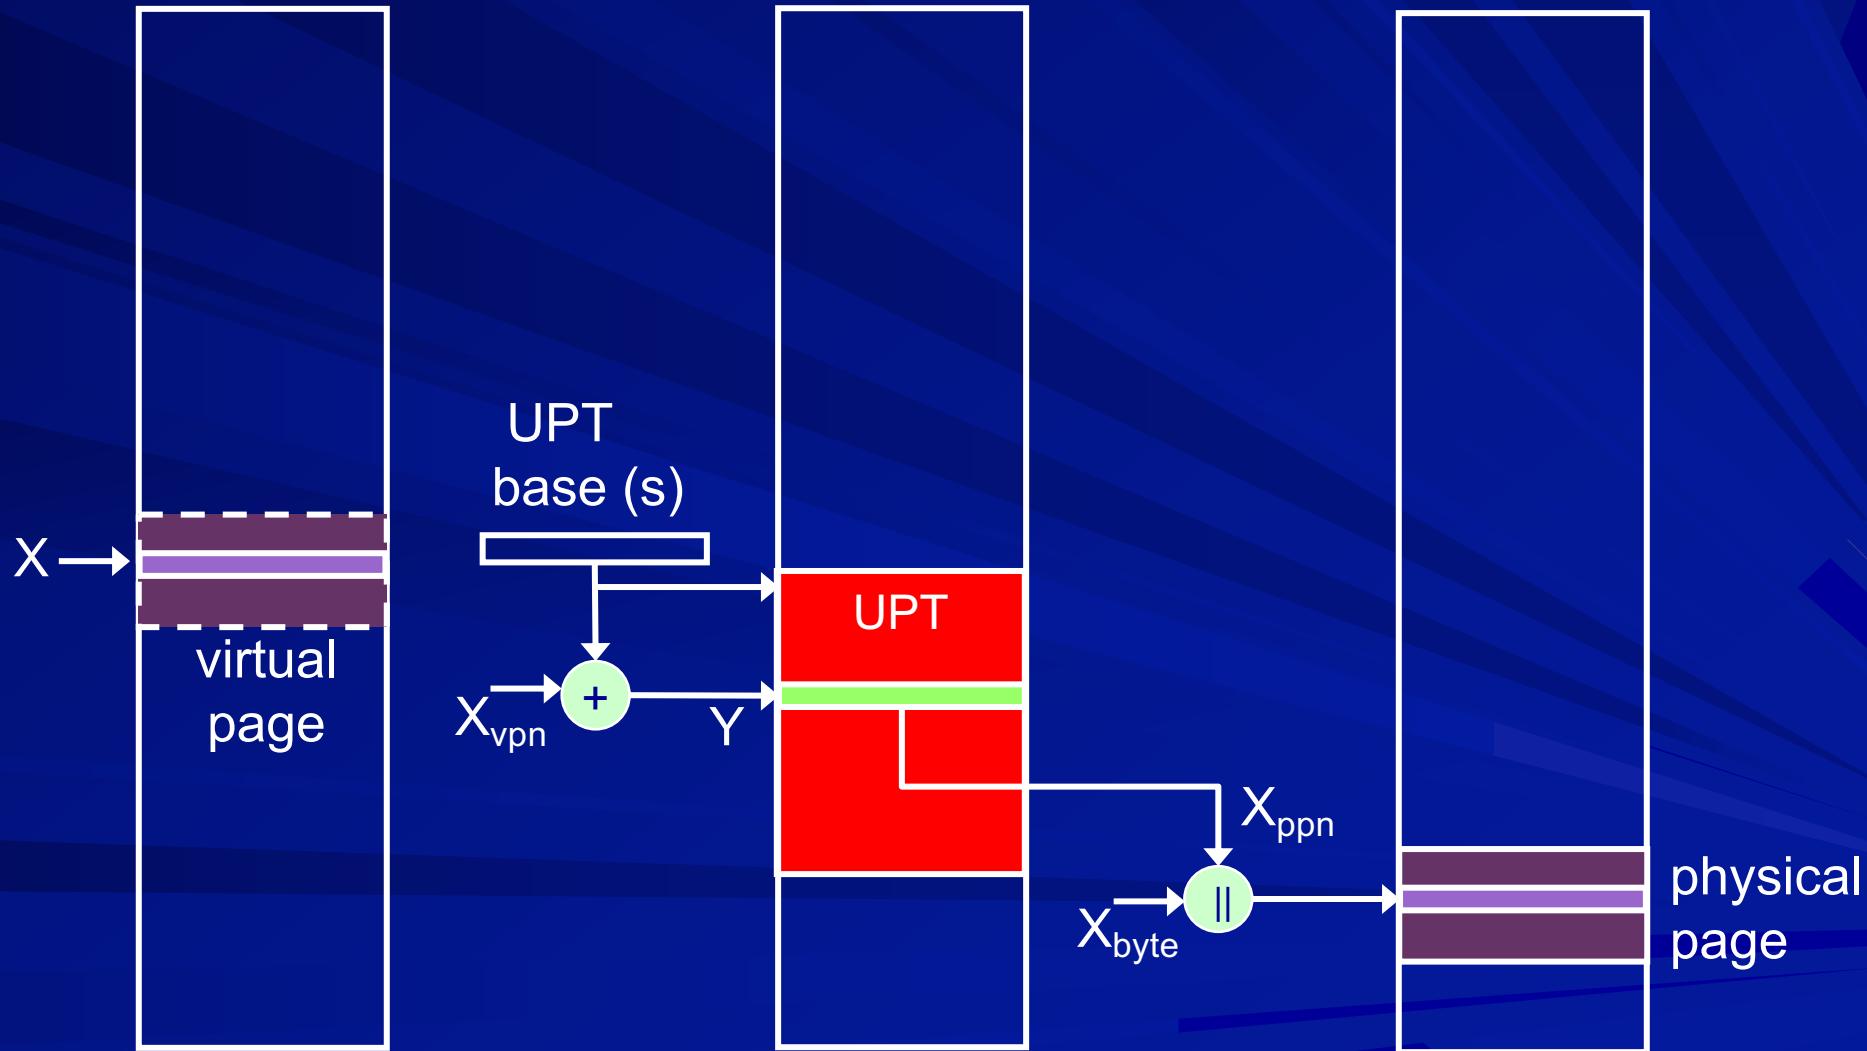
physical
space



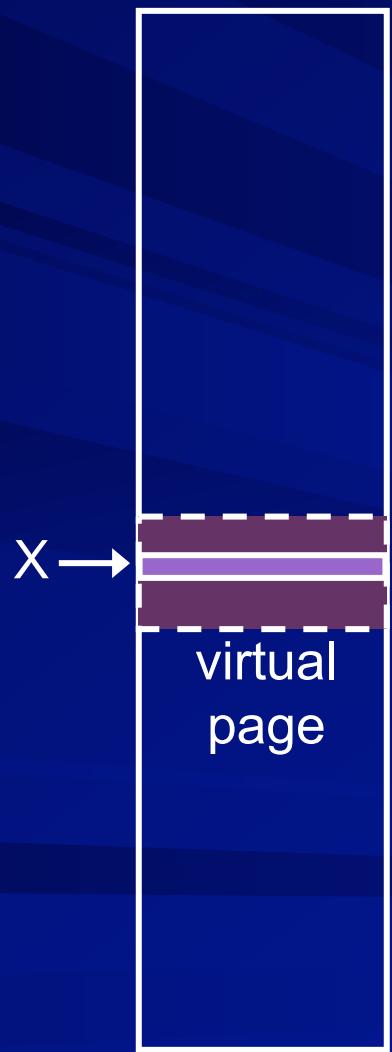
user
virtual space

system
virtual space

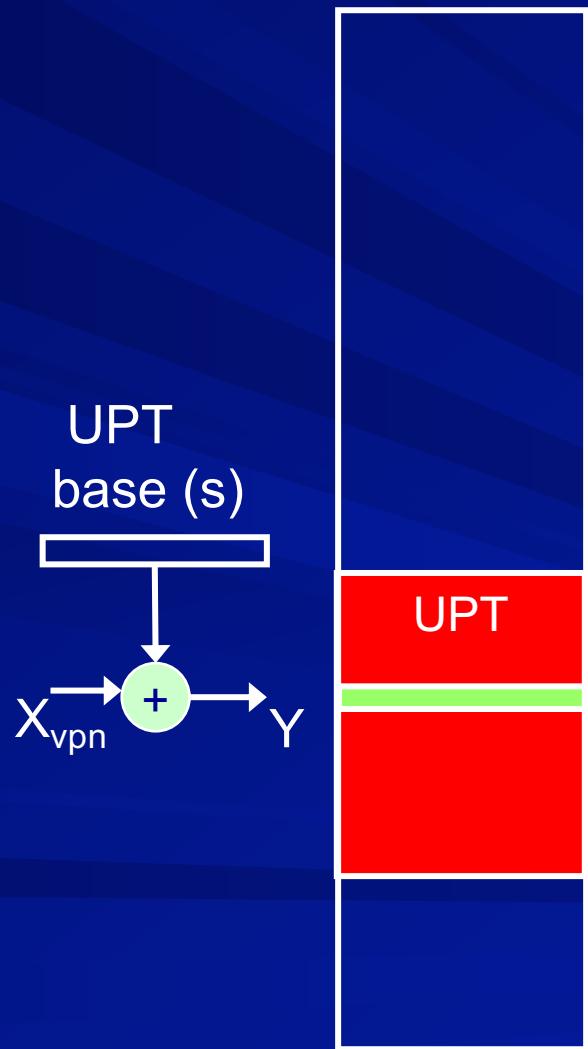
physical
space



user
virtual space



system
virtual space



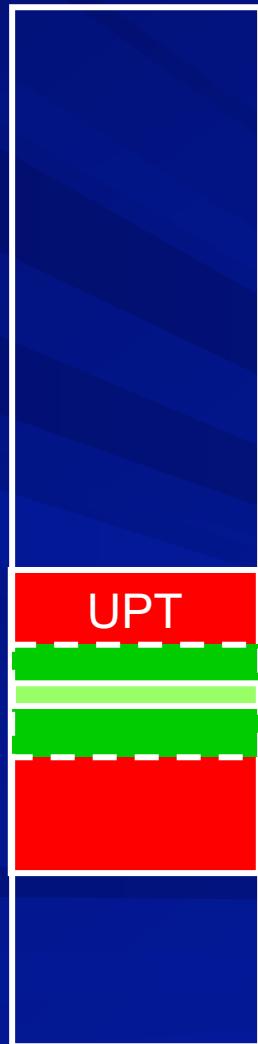
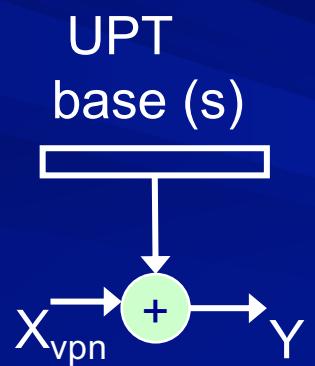
physical
space



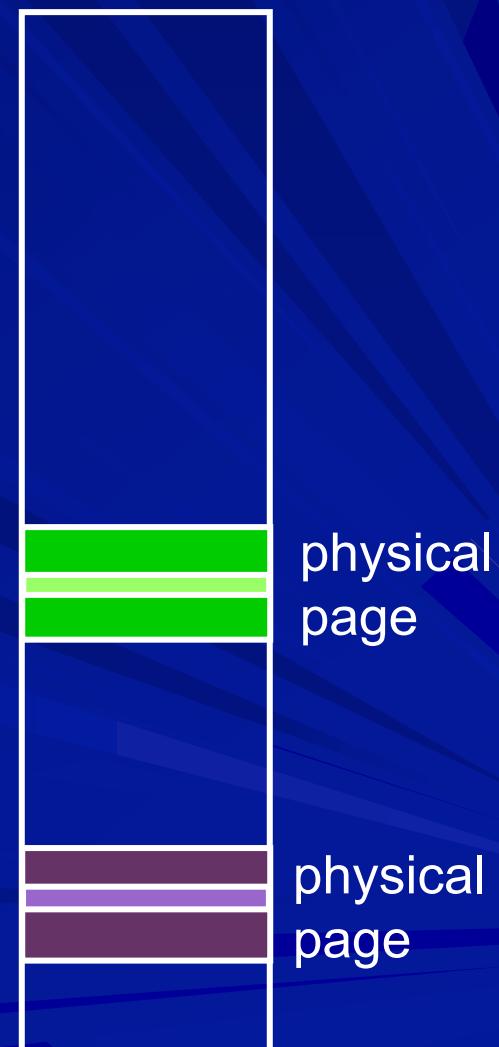
user
virtual space



system
virtual space



physical
space



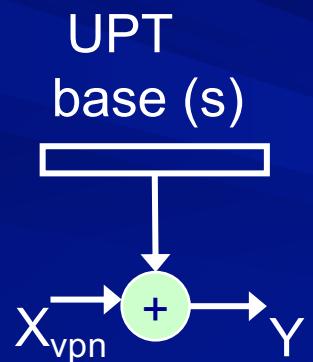
physical
page

physical
page

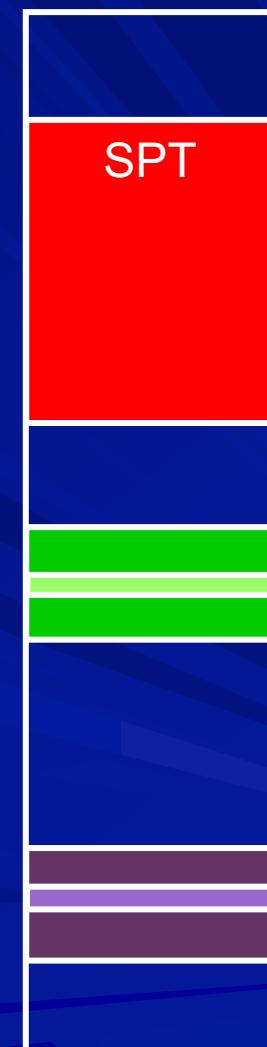
user
virtual space



system
virtual space



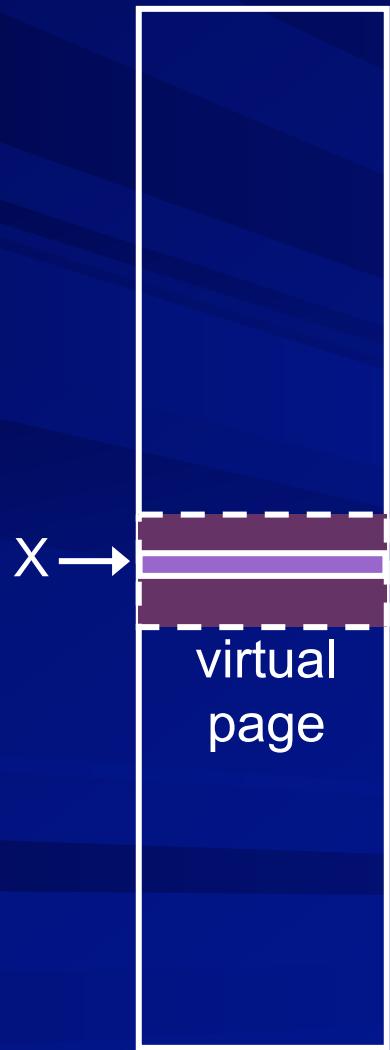
physical
space



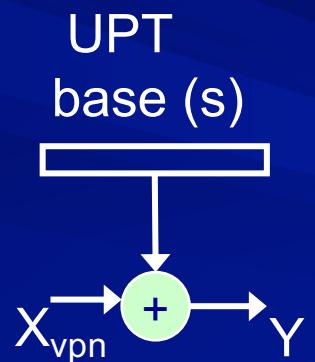
physical
page

physical
page

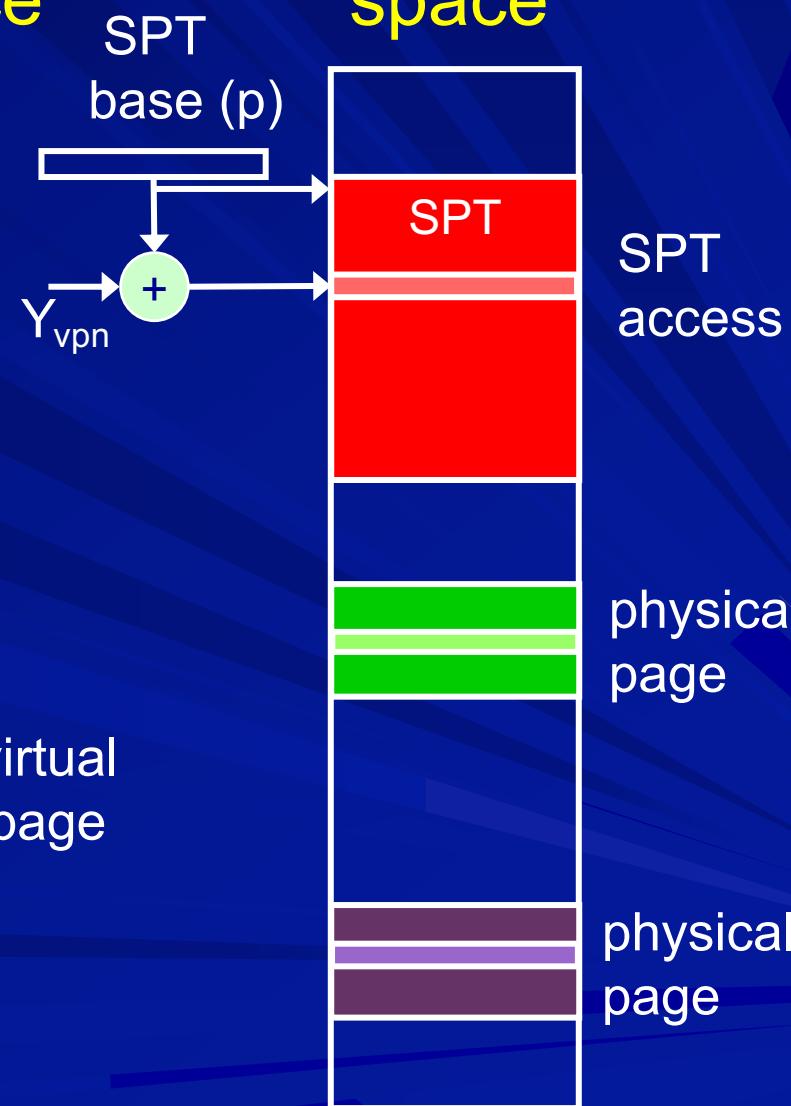
user virtual space



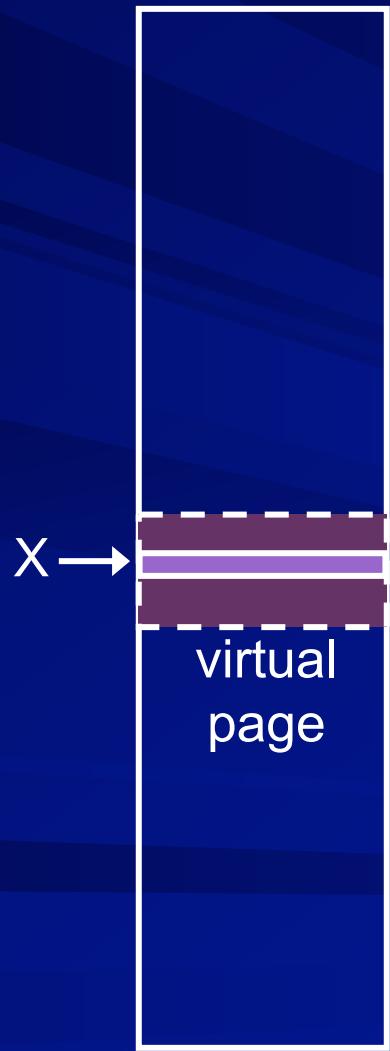
system virtual space



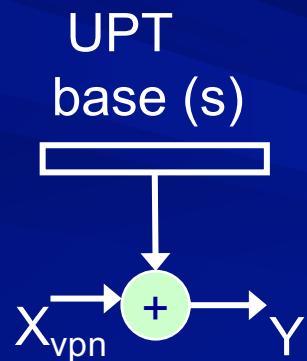
physical space



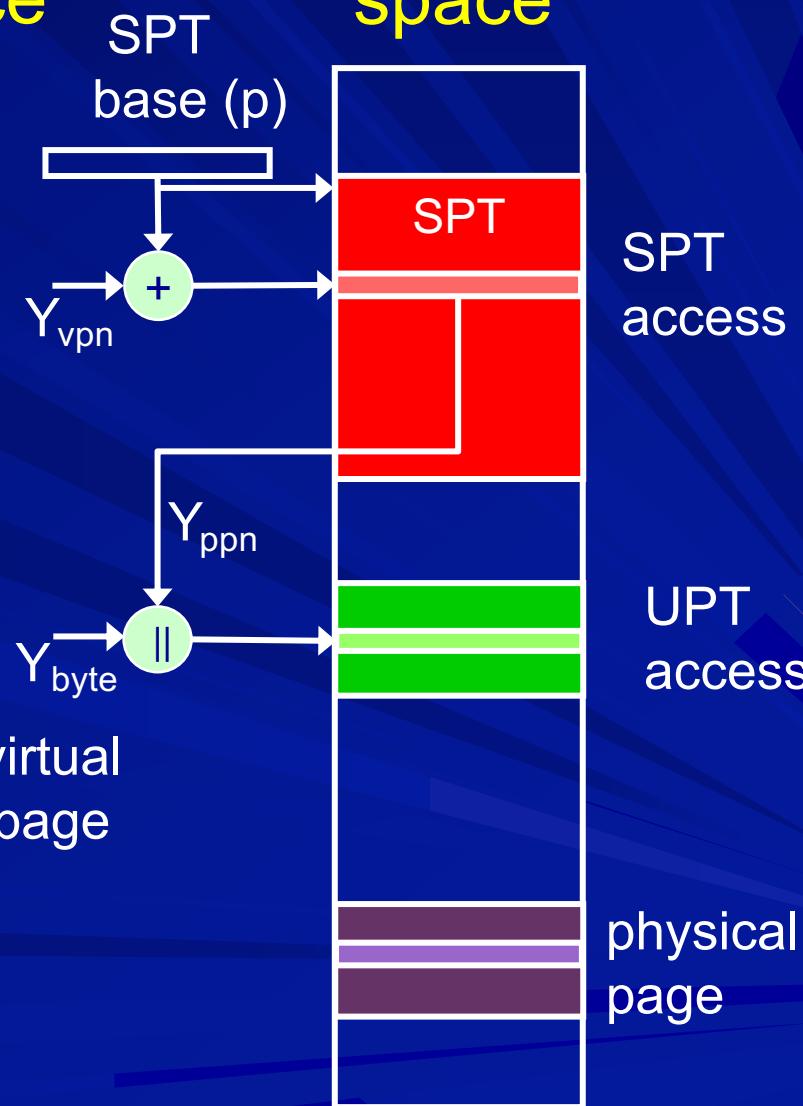
user virtual space



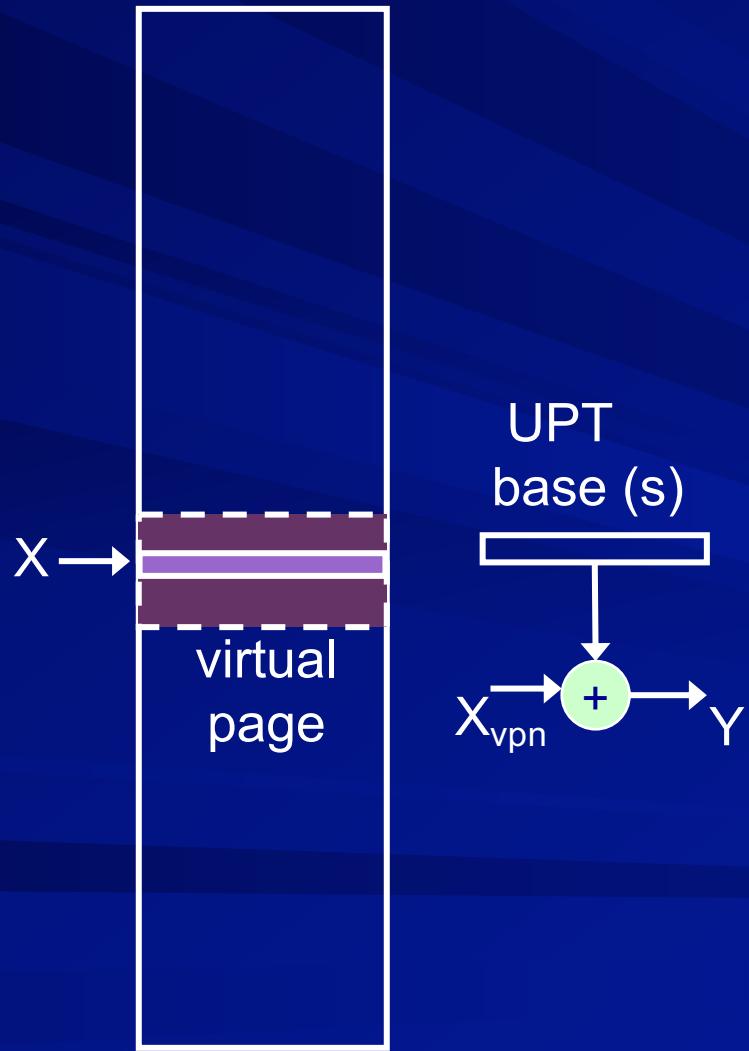
system virtual space



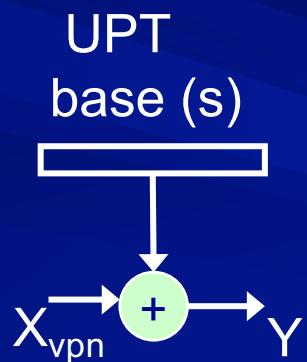
physical space



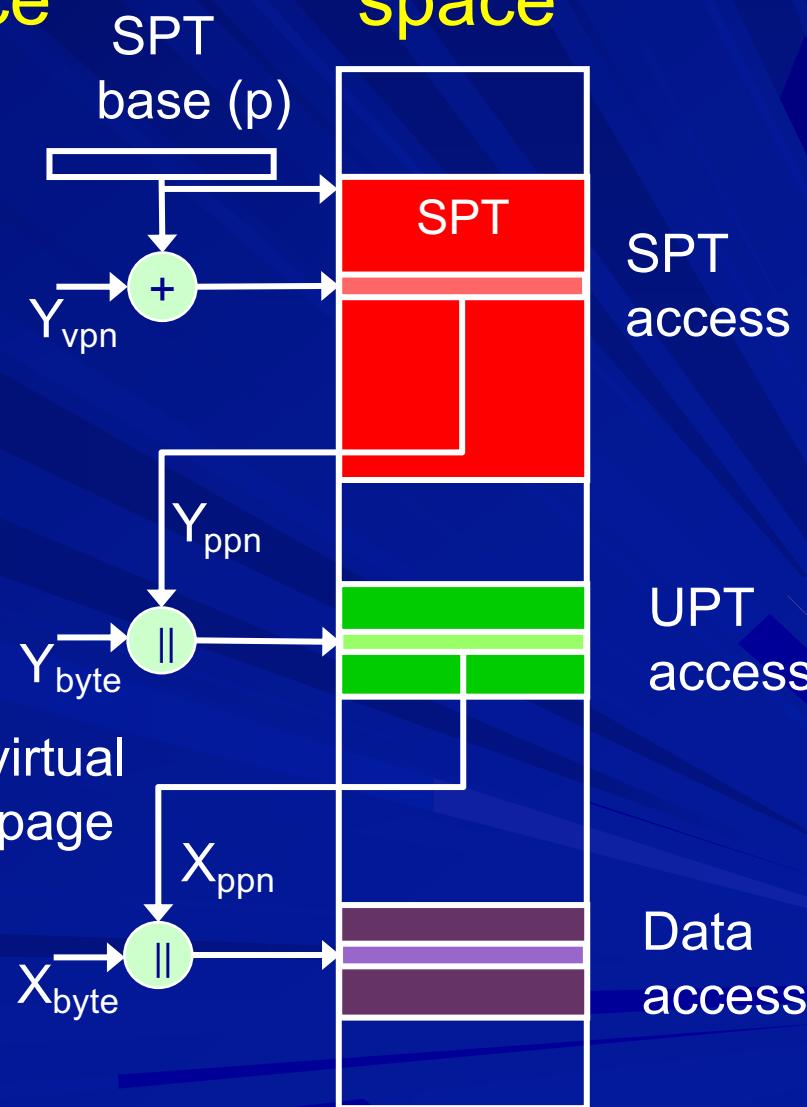
user virtual space



system virtual space



physical space



SPT
access

UPT
access

Data
access

Memory protection with VM

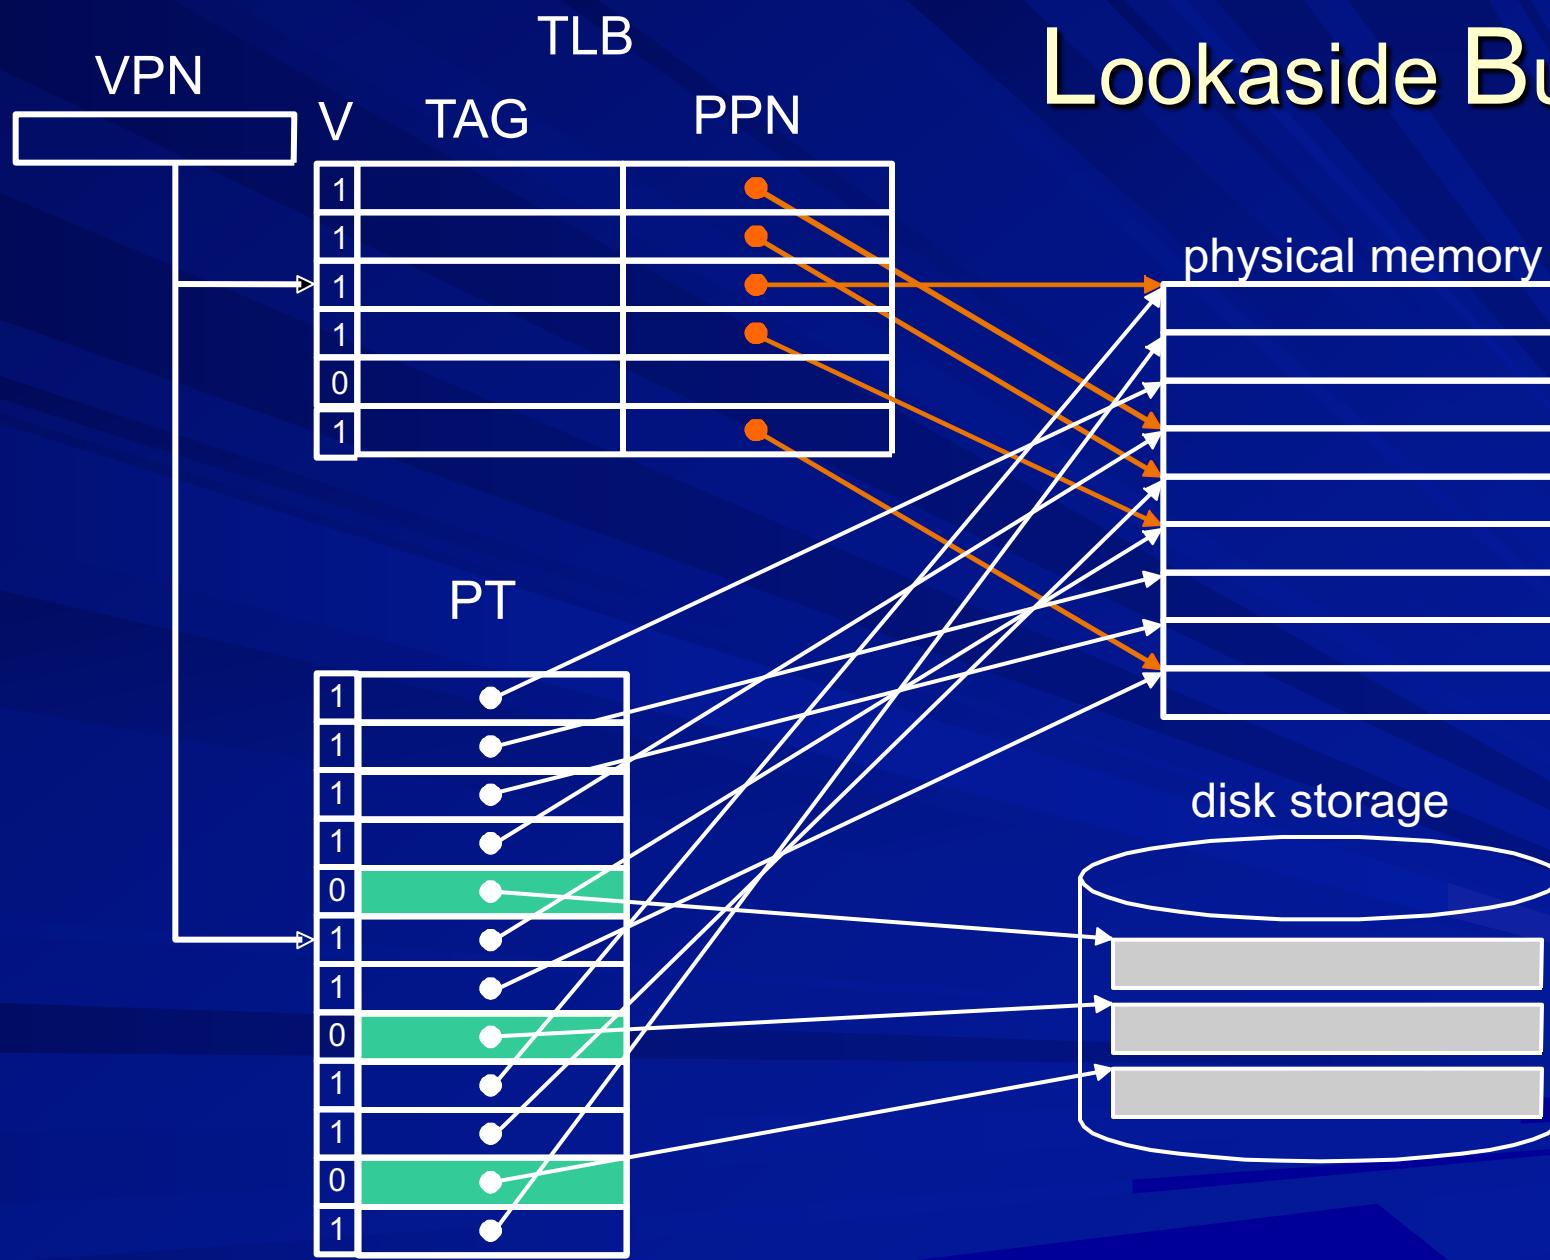
- A process can access only those physical pages which are pointed to by its page table
- Page tables are initialized by the OS and are updated automatically. User processes can't modify these.
- Programs run in user mode or privileged mode. Certain operations are possible in privileged mode only.

Virtual memory hit time

- simple page table :
 $2 \times$ physical memory access time
- 2 level page table :
 $3 \times$ physical memory access time
- paged page table :
 $3 \times$ physical memory access time

Can this be better??

Translation Lookaside Buffer



How is TLB Miss handled?

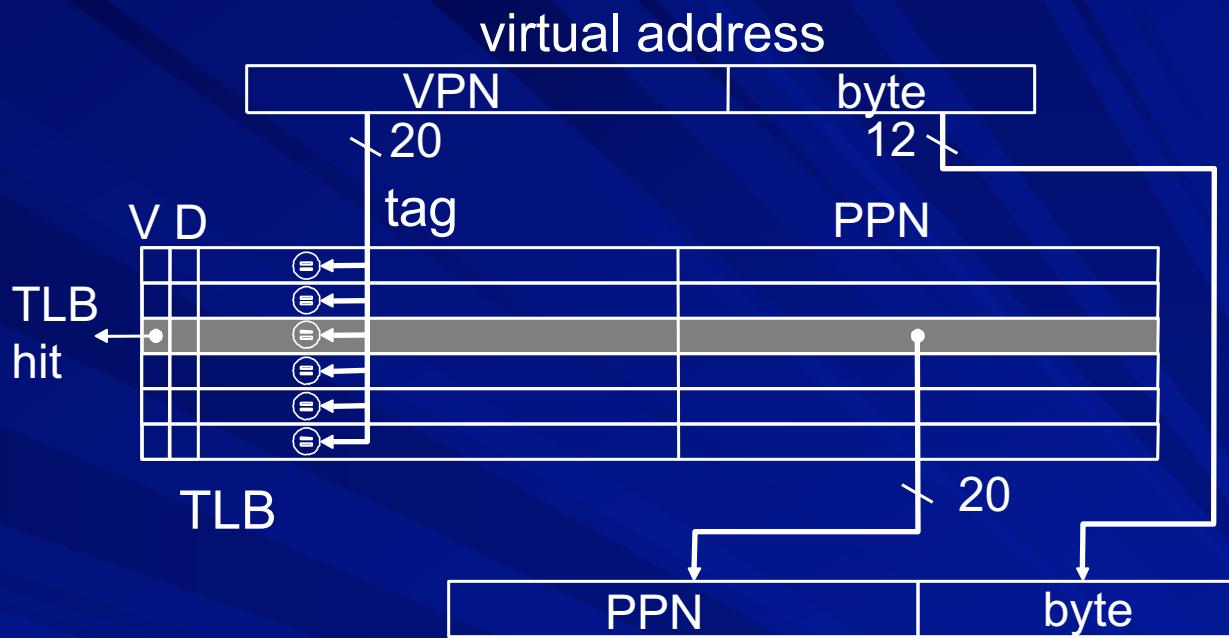
■ Hardware

- Memory Management Unit accesses Page Table
- Extra hardware, but efficient

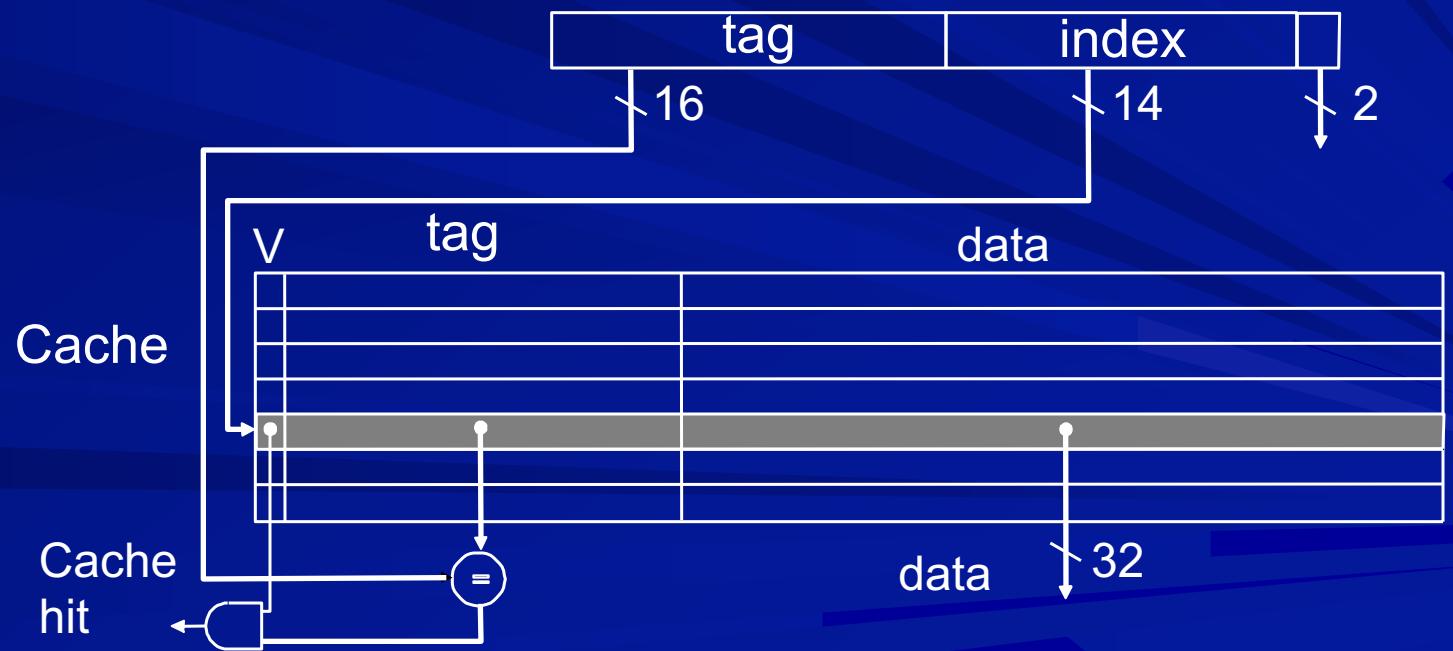
■ Software

- Exception is caused
- OS (exception handler) accesses Page Table
- Page Table structure defined by OS, not fixed by hardware

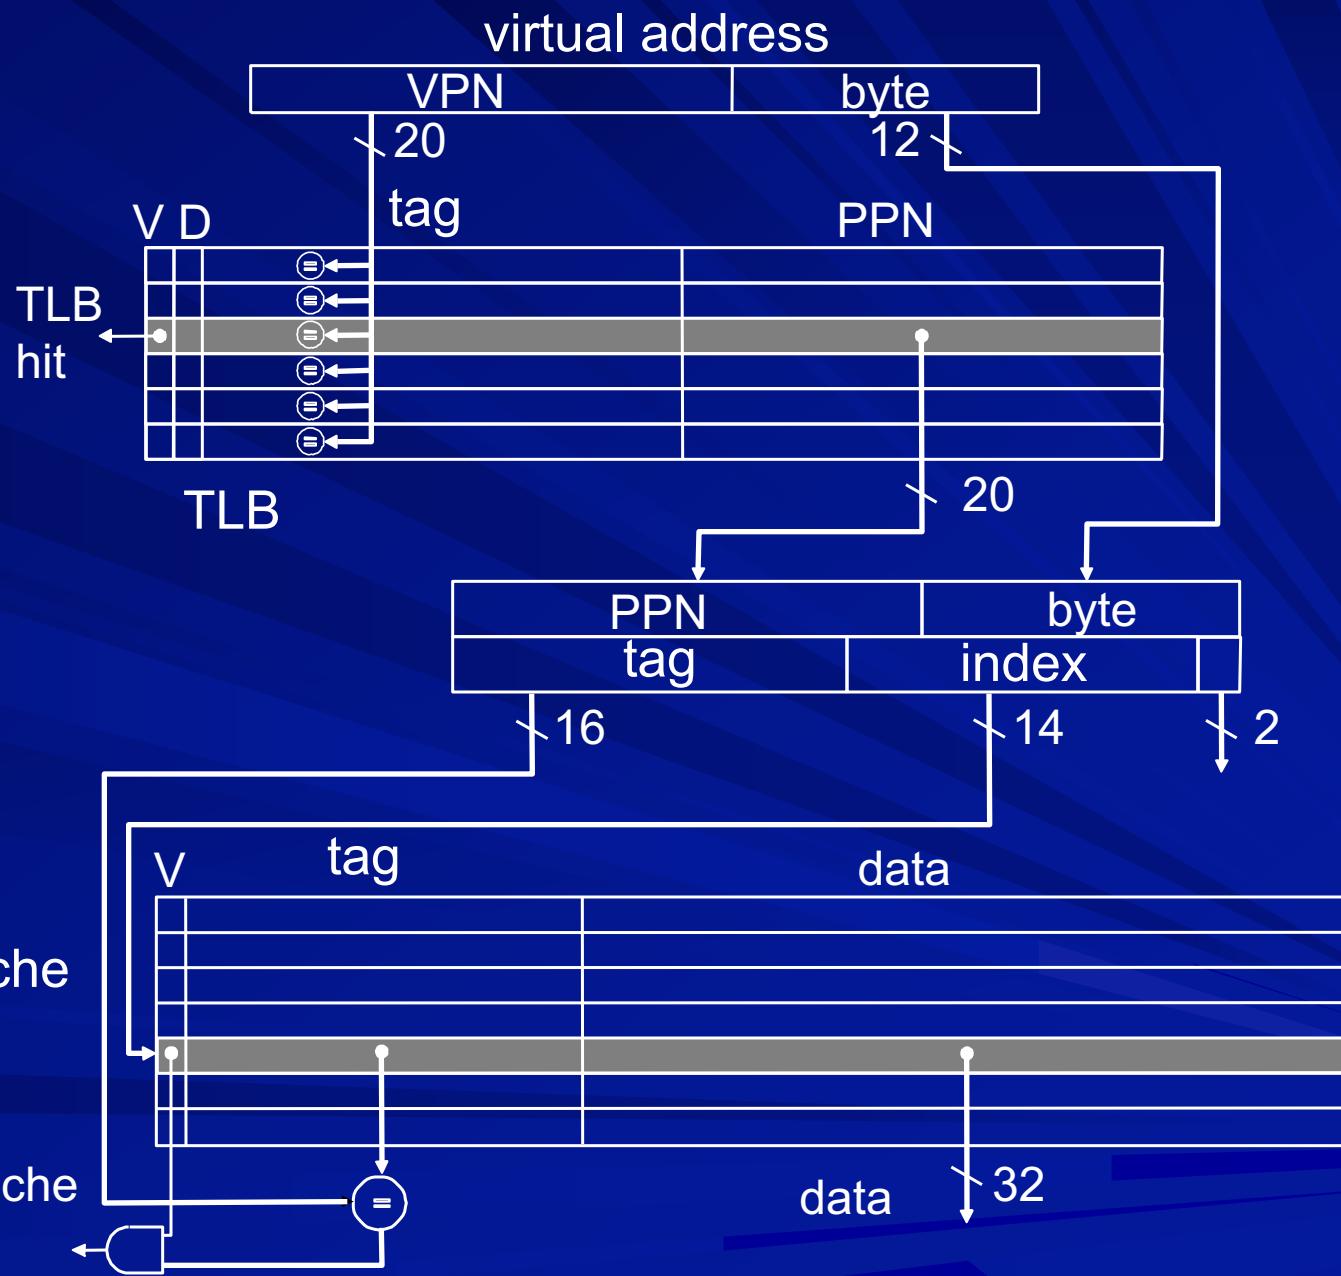
TLB with Cache



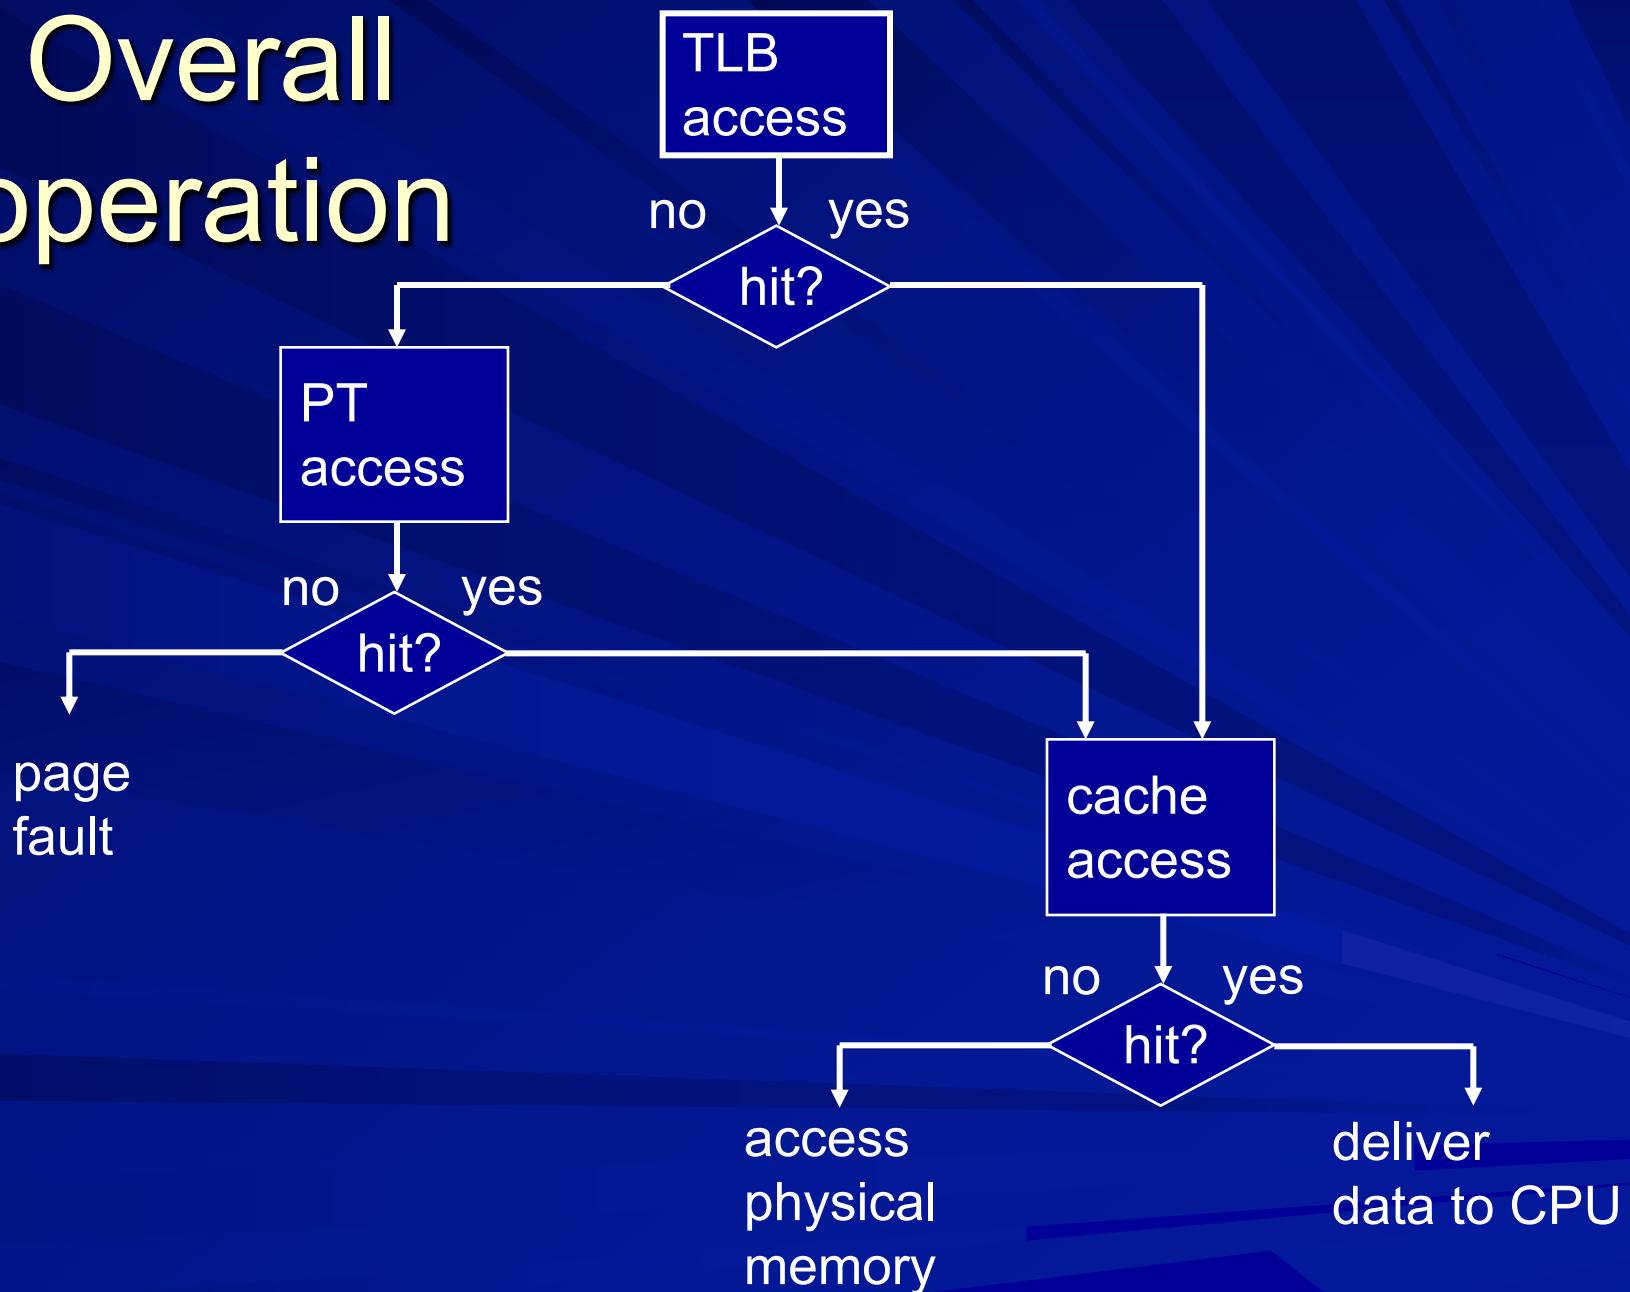
TLB with Cache



TLB with Cache



Overall operation



Virtually addressed cache

- Tags in the cache are from virtual addresses
- Cache access is made first
- TLB accessed only after cache miss
- Problem of aliasing - two copies of a shared object in physical memory
- Physically tagged and virtually indexed cache may be used

Indexing before address translation

- Virtually indexed, physically tagged cache
 - simple and effective approach
 - indexing overlaps with address translation, tag matching after address translation
 - valid only if index field does not change during address translation (possible for caches that are not too large)

Virtual Page No.	Offset in page
Tag	Index